

Silicon Graphics, Inc.

XFS Overview & Internals

15 - DAPI

© Copyright 2006 Silicon Graphics Inc. All rights reserved.

Permission is granted to copy, distribute, and/or modify this document under the terms of the Creative Commons Attribution-Share Alike, Version 3.0 or any later version published by the Creative Commons Corp. A copy of the license is available at <http://creativecommons.org/licenses/by-sa/3.0/us/>.

November 2006

DMAPI

- Data Management Applications Programming Interface
- Developed by Data Management Interfaces Group consortium (DMIG), 1993-1996
- Provides standardized access to filesystems for Hierarchical Storage Managers (HSMs) and backup packages
- DAPI in SLES is used by vendors other than just SGI
 - IBM
 - HP

HSM Example

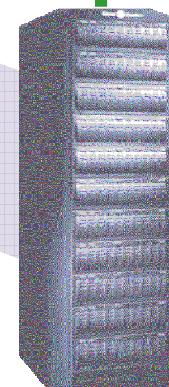
Primary Storage

Online high-performance disk
e.g. 2 to 10 hours
current projects



Nearline Disk

High Capacity, Low cost,
Lower performance
e.g. 6 to 20 hours
this years projects



Tape Library

Archive
e.g. hundreds to
thousands of hours



*Promote
on demand*

*Promote
on demand*

*Demote
e.g. > 30 days*

*Demote
e.g. > 12 months*

HSM Economic Benefit

Example Environment

2TB data less than 30 days old

6TB data between 30 days and one year old

30TB data over one year old

Without HSM

High-Perf. Disk
30TB



With HSM

High-Perf. Disk
2TB



Low Cost
Disk
6TB



Tape Library
30TB



<u>System</u>	<u>c/MB</u>	<u>w/o DMF</u>	<u>With DMF</u>
High-Perf. Disk	4c	\$900K	\$80K
Low Cost Disk	0.8c	-	\$50K
Tape Library	0.37c	-	\$110K
DMF		-	\$110K
Total Investment		\$900K	\$350K

What do HSMs need?

- To be notified of accesses/changes to a file's metadata
- To be able to accept or reject such accesses and changes
- To have user processes block until the change has been approved
- To be able to access a file without knowing a pathname
- To be notified of changes to a file's metadata
- To be able to read and change the contents of a file without blocking
- To save HSM-specific information persistently with a file
- To scan all files in a filesystem quickly
- To be notified when filesystems become full

DMAPI Standard

- DAPI Versions 2.1 and 2.3
- X/Open Data Storage Management API (XDSM) - 2/97
- Not a rigorous standard.
 - ~70 function calls
 - ~40 structures
 - 24 events
- Spec available free in HTML format from
<http://www.opengroup.org/pubs/catalog/c429.htm>
- Alex Miroshnichenko's paper from AUUG 1996:
<http://www.csu.edu.au/special/auugwww96/proceedings/alex/alex.html>

Handles

- Allow file access without pathnames
 - Similar to NFS file handles
- A handle is an opaque persistent identifier which is unique per host
 - Should not change for the life of the object
- Handles are represented as variable length byte streams
 - `void*` and `size_t`
- An existence of a valid handle does not guarantee existence of the object referenced by it.
 - The only way to guarantee the existence of the object referred to by a handle is to either obtain a right or hold on the object
- Three handle types:
 - Global handle - (used only in mount requests)
 - Filesystem handle - (fsid hash)
 - File handle - (fsid hash/fid_t structure)

Handles - DMAPI Functions

- dm_path_to_handle
 - create file handle from pathname
- dm_fd_to_handle
 - create file handle from file descriptor
- dm_path_to_fhandle
 - create filesystem handle from file descriptor
- dm_handle_to_fhandle
 - extract the filesystem handle from a file handle
- dm_handle_cmp
 - file handle comparison
- dm_handle_is_valid
 - determine if a handle is valid
- dm_handle_hash
 - hash contents of a handle
- dm_handle_free
 - free storage allocated to a handle
- dm_handle_to_path
 - get a pathname from two handles

DMAPI Events and Tokens

- Events are messages from the kernel to a DAPI application
- Provide notification of accesses/changes to a file or filesystem
- When a synchronous event is generated, the user process is suspended in the kernel until a DAPI application issues an explicit response to the event
 - A token is a reference to the kernel state associated with a synchronous event
- Asynchronous events are for notification purpose only
 - May indicate a completion (or failure) of certain operations
 - Do not block user processes
- All contain handles, directory names, etc.

Events - Filesystem Admin

- mount
 - a filesystem is about to be mounted
- preunmount
 - a filesystem is about to unmount
- unmount
 - an unmount succeeded/failed
- nospace
 - filesystem has run out of space
- Not implemented by XFS:
 - debut
 - inode has been read from disk

Events - Namespace

- **create/postcreate**
 - create a file or directory
- **remove/postremove**
 - unlink a filesystem object
- **rename/postrename**
 - rename a filesystem object
- **link/postlink**
 - (hard) link a filesystem object
- **symlink/postsymlink**
 - create a symlink

Events - Data

- read
 - the specified file byte range is about to be read
- write
 - the specified file byte range is about to be written
- truncate
 - file is about to be truncated to the specified address (away from or towards zero)

Events - Metadata

- attribute
 - a filesystem object's attribute has changed
- destroy
 - a filesystem object has been destroyed
- Not implemented by XFS:
 - cancel
 - a previously issued event is cancelled
 - close
 - a filesystem object has been closed

Events - Pseudo

- user event
 - a message between cooperating DMAPI processes

Sessions and Dispositions

- DMAPI uses sessions as the primary communication channels between DM application and the kernel component of DMAPI
 - Almost all DMAPI calls require a session argument
- Sessions and event lists allow DM application to exercise a fine control of the event delivery and generations
 - Applications may direct different event types to different sessions and avoid session overloading
- Dispositions indicate which event types should be delivered to this session.
 - A DM application creates a session by calling `dm_open_session()` and then may register event dispositions for this session by calling `dm_set_disp()`

Connecting Events with Sessions

- Select desired events to be generated
 - dm_eventlist_t bitmap
- Set by filesystem
 - not persistent
- Set by file
 - persistent
- Set globally
 - to trigger mount events
- Next set disposition of those events
 - recipient session

Connecting Events - DMAPI Functions

- dm_get_eventlist
 - get list of enabled events for object
- dm_set_eventlist
 - set list of enabled events for object
- dm_set_disp
 - set disposition of events on a filesystem to a particular session
- dm_set_return_on_destroy
 - specify a DMAPI attribute to return with destroy events

Managing Sessions - DMAPI functions

- dm_create_session
 - create a new session
- dm_query_session
 - query a session for information
- dm_destroy_session
 - destroy the specified session
- dm_getall_sessions
 - get all extant sessions

Processing Events - DMAPI Functions

- dm_get_events
 - get next available event messages
- dm_respond-event
 - respond to one event
- dm_pending
 - notify FS of slow application event
- dm_create_userevent
 - generate a user pseudo-event
- dm_send_msg
 - send message to indicated session
- dm_move_event
 - move event to another session

Session Recovery

- dm_getall_sessions
 - get all extant sessions
- dm_query_session
 - query a session for info
- dm_create_session
 - create/assume a session
- dm_getall_tokens
 - get outstanding tokens for a session
- dm_find_eventmsg
 - get message for an event
- dm_getall_disp
 - get event dispositions for all filesystems for a session

Managed Regions

- A managed regions is the mechanism for an application to control file data access at a granularity level less than file size. It is an extent in the logical file space
 - starting offset
 - length
 - event generation flags
- Provide byte-granularity events for all accesses/modifications of a file's data
 - Only apply to file objects
- Cover non-overlapping byte ranges in a file
- Separately detect reads, writes, truncates
- Only one, persistent in XFS

Managed Regions – DMAPI Functions

- dm_get_region
 - get managed regions for a file
- dm_set_region
 - set managed regions for a file

File Access and Modification

- Must not change mtime or atime (ctime changes)
 - TODO: Bug
- Must not block
- Must not generate events

File Access and Modification – DMAPI Functions

- dm_probe_hole
 - return rounded result of the area where a hole is to be punched
- dm_punch_hole
 - create a hole in a file
- dm_read_invis
 - read a file bypassing DMAPI events
- dm_write_invis
 - write a file bypassing events
- dm_sync_by_handle
 - sync a file's data to disk
- dm_get_fileattr
 - return file attributes
- dm_set_fileattr
 - set file time stamps, ownership, mode, length, etc.
- dm_get_allocinfo
 - return allocation info for a file

Persistent DMAPI Data

- Stored per file as SGI_DMI_xxx extended attribute
- Persistent across reboots
- DMF stores in SGI_DMI_DMATTR
 - bfid, file state
 - version info
 - Flags
- inodes should be 512 bytes (-i size=512)

Persistent DMAPI Data – DMAPI Functions

- dm_get_dmattr
 - retrieve a DMAPI attribute
- dm_set_dmattr
 - set/replace a DMAPI attribute
- dm_remove_dmattr
 - remove a DMAPI attribute
- dm_getall_dmattr
 - retrieve all DMAPI attributes for a file

Scanning Filesystems Quickly

- dm_init_attrloc
 - initialize a bulk attribute location offset
- dm_get_bulkattr
 - get bulk attributes for entire filesystem
- Not implemented by XFS:
 - dm_get_bulkall
 - get data management attributes for entire filesystem

Rights

- DM applications use the tokens to obtain access rights to file system objects to guarantee stability of the objects.
 - Allow session to control access/updates
 - Provide for atomic updates over many functions
- Access rights may be *shared* and *exclusive*
 - Shared right allows read-only access (DM_RIGHT_SHARED)
 - Exclusive right allows modification to the object (DM_RIGHT_NULL)
- Pseudo-available in XFS

Rights – DMAPI Functions

- dm_release_right
 - release all access rights to an object
- dm_query_right
 - determine set of access rights to an object
- dm_upgrade_right
 - upgrade current right to exclusive
- dm_downgrade_right
 - downgrade current right to shared
- Not implemented in XFS:
 - dm_request_right
 - request a specific access right

XFS DMAPI Mount Options

- Options must be used together
- dmapi
 - Enable the DMAP event callouts
 - xdsm and dmi options are equivalent
- mtpt=mountpoint
 - The mountpoint specified here will be included in the DMAP mount event, and should be the path of the actual mountpoint that is used.

```
> grep dmapi /etc/fstab  
/dev/sdb1 /mnt/scratch1 xfs dmapi,mtpt=/mnt/scratch1,defaults 0 0
```

```
> mount  
...  
/dev/sdb1 on /mnt/scratch1 type xfs (rw,dmapi,mtpt=/mnt/scratch1)  
...
```

What if there is no DMAPI module?

```
# mount -o logdev=/dev/sdb,dmapi,mtpt=/mnt/data /dev/sdd /mnt/data
mount: wrong fs type, bad option, bad superblock on /dev/sdd,
      missing codepage or other error
      In some cases useful info is found in syslog - try
      dmesg | tail  or so
# dmesg | tail | grep dmapi
XFS: unknown mount option [dmapi].
```

- Check that DMAPI is probed on boot and loaded
 - SLES10 set DMAPI_PROBE to yes in /etc/sysconfig/sysctl
- Without rebooting XFS with load dmapi with

```
# echo 1 > /proc/sys/fs/xfs/probe_dmapi
# lsmod | grep dmapi
xfs_dmapi          22936  1
dmapi              39728  1 xfs_dmapi, [permanent]
xfs                479796  4 xfs_dmapi, xfs_quota
```

DMAPI Implementation

- Not all DMAPI interfaces used by SGI's HSM product
 - Many of the bugs recently fixed in DMAPI did not impact SGI's HSM
- Some DMAPI interfaces are not implemented by XFS
- Maintaining 2.4/2.6 compatible code makes it difficult to read
 - Plan is to remove 2.4 support from this code base
- DMAPI requires two kernel changes that are not in mainline
 - mprotect
 - open_exec

DMAPI Libraries

```
> rpm -qpl dmapi-2.2.3-12.2.ia64.rpm
/lib/libdm.so.0
/lib/libdm.so.0.0.4
/usr/share/doc/packages/dmapi
/usr/share/doc/packages/dmapi/CHANGES.gz
/usr/share/doc/packages/dmapi/COPYING
/usr/share/doc/packages/dmapi/PORTING
/usr/share/doc/packages/dmapi/README
```

DMAPI Development Package

```
> rpm -qpl dmapi-devel-2.2.3-12.2.ia64.rpm
/lib/libdm.so
/usr/include/xfs/dmapi.h
/usr/lib/libdm.a
/usr/lib/libdm.la
/usr/lib/libdm.so
/usr/share/man/man3/dmapi.3.gz
```

DMAPI Kernel Modules

```
> rpm -qpl kernel-default-2.6.16.21-0.8.ia64.rpm | grep dmapi
/lib/modules/2.6.16.21-0.8-default/kernel/fs/dmapi
/lib/modules/2.6.16.21-0.8-default/kernel/fs/dmapi/dmapi.ko
  – Filesystem independent module
/lib/modules/2.6.16.21-0.8-default/kernel/fs/xfs/dmapi
/lib/modules/2.6.16.21-0.8-default/kernel/fs/xfs/dmapi/xfs_dmapi.ko
  – XFS dependent module
```

DMAPI Triage

